



Institute of Electronics
National Yang Ming Chiao Tung University
Hsinchu, Taiwan

AI Training Course Series

Model Compression: Quantization + Pruning + Weight-Sharing

Lecture 12



Student: Ting-Wei Hu
Advisor: Juinn-Dar Huang, Ph.D.

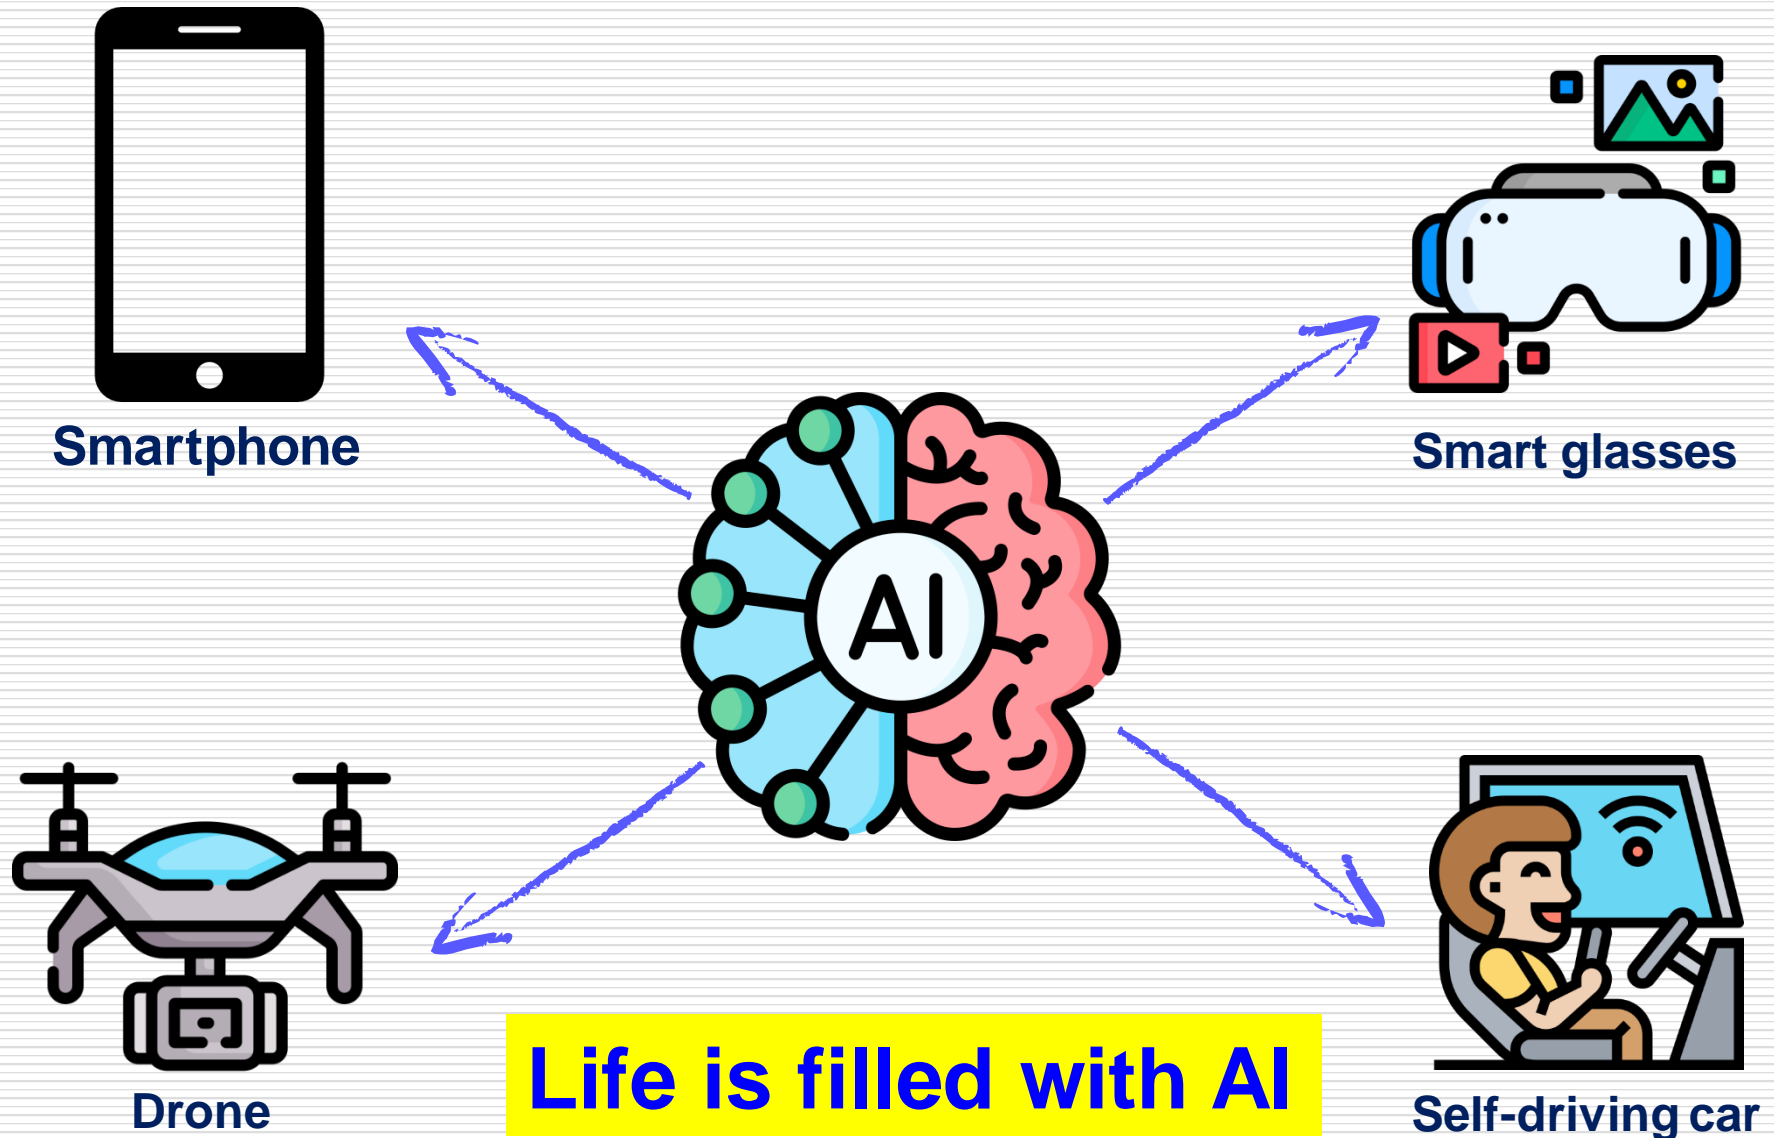
Aug 22, 2024

Outline

- Necessity of model compression
- Pruning
- Weight-sharing
- Low-rank approximation
- Quantization concepts
- Uniform integer quantization
- Survey of LLM quantization
- PFPQuant
- Homework


Necessity of Model Compression

AI Applications (1/2)





AI Applications (2/2)

- Requirements for the common applications
 - Low latency
 - Energy efficient



Self-driving car

high latency ☹️




Drone

high energy consumption ☹️




Booming AI Models

- Recent deep learning models are **complex**
 - **LeNet-5**: 1M parameters (1998)
 - **VGG-16**: 133M parameters (2014)
 - **CoCa**: 2100M parameters (2022)
- 83× ↓ **Emergence of LLMs**
- **ChatGPT**: **175B** parameters (2023) **extremely huge!**



Computational Issue (1/2)

- The model size is developing at a faster pace than the GPU memory in recent years




Computational Issue (2/2)

- LLMs are compute and memory-intensive
- Swin Transformer model in computer vision → Swin-L **197M** parameters
- ChatGPT → **175B** parameters consuming at least 350GB memory to store in FP16
 - 5 X 80GB A100 GPUs for inference
 - Huge computation and communication overhead
 - Almost **1000×** increase compared with Swin-L

Energy Issue

- DRAM results in significant energy consumption
- Huge models usually access data from DRAM
 - More data transfer, more energy consumption

Operation	Energy [pJ]	Relative Cost
32 bit int ADD	0.1	1
32 bit float ADD	0.9	9
32 bit Register File	1	10
32 bit int MULT	3.1	31
32 bit float MULT	3.7	37
32 bit SRAM Cache	5	50
32 bit DRAM Memory	640	6400



Energy consumption increases dramatically

Why Model Compression


- Challenges of complex deep models
 - Huge storage (memory, disk) requirement
 - Computationally expensive
 - Consume lots of energy
 - Hard to deploy models on small edge devices

Compression becomes more significant !

Pruning


Introduction

- Motivated by how real brain learns
 - Usefulness in, waste out
- Propose **pruning** technique discarding relatively unimportant weights to achieve sparse model
 - Remove redundant connections




How Pruning Works (1/2)

- Threshold pruning
 - Set a parameter as threshold
 - Remove weights whose $\text{abs}(\text{weight}) < \text{threshold}$



How Pruning Works (2/2)

- Retrain after threshold pruning
 - Reduce accuracy drop
 - Learn effective connections by **iterative pruning**



**repeat the two steps iteratively
To achieve gradual pruning**

Retraining Strategies

- After pruning, the **performance loss** should be **compensated by retraining**
 - **Prune weights of all layers once** then retrain few epochs, then repeat this procedure until certain degree of accuracy is restored
 - › Faster
 - › Lower accuracy
 - **Prune weights layer by layer** then retrain iteratively, the model is retrained before pruning the next layer
 - › Slower
 - › Higher accuracy


Model Compression Rate

- Result of threshold pruning
 - Experimented with LeNet, AlexNet, VGG (ImageNet)
 - About **10x smaller** compared to origin network
 - Mostly contributed by **FC layers**
 - Less than 1% accuracy loss


Network	Top-1 Error	Top-5 Error	Parameters	Compression Rate
LeNet-300-100 Ref	1.64%	-	267K	
LeNet-300-100 Pruned	1.59% 3.0%↓	-	22K	12×
LeNet-5 Ref	0.80%	-	431K	
LeNet-5 Pruned	0.77% 3.7%↓	-	36K	12×
AlexNet Ref	42.78%	19.73%	61M	
AlexNet Pruned	42.77% 0.2%↓	19.67%	6.7M	9×
VGG-16 Ref	31.50%	11.32%	138M	
VGG-16 Pruned	31.34% 5.0%↓	10.88%	10.3M	13×

Unstructured Pruning

- Unstructured pruning
 - Better accuracy
 - Needs extra memory to store index
 - Extremely hard to keep parallelism in hardware acceleration (i.e., you'd better forget about it!)



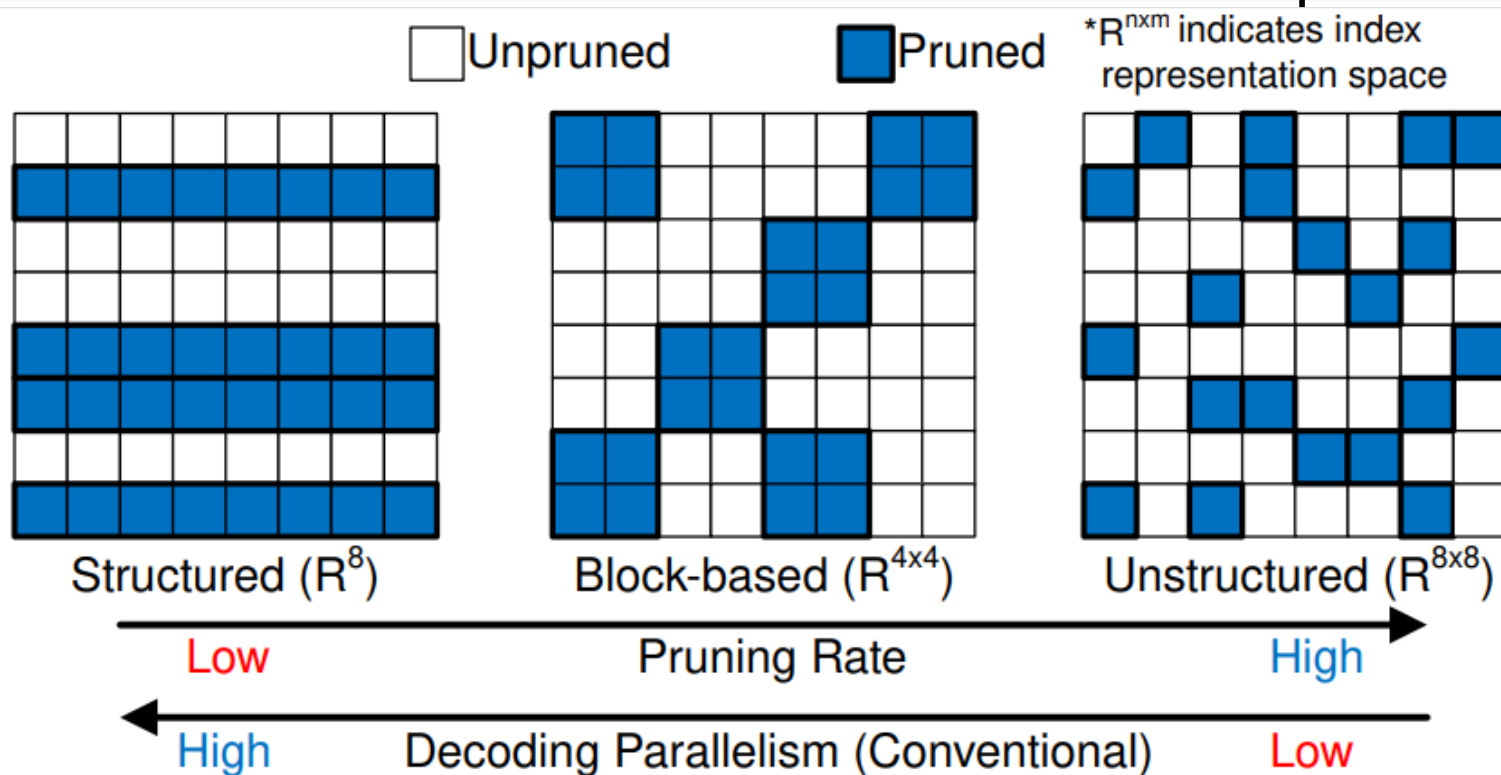
unstructured



structured

Structured Pruning


- Structured pruning
 - Maintains the regularity of the weight matrix
 - **Accuracy degradation** due to larger information loss
 - **Eliminates the overhead** and accelerates computation



Weight-Sharing

Weight-Sharing By Clustering

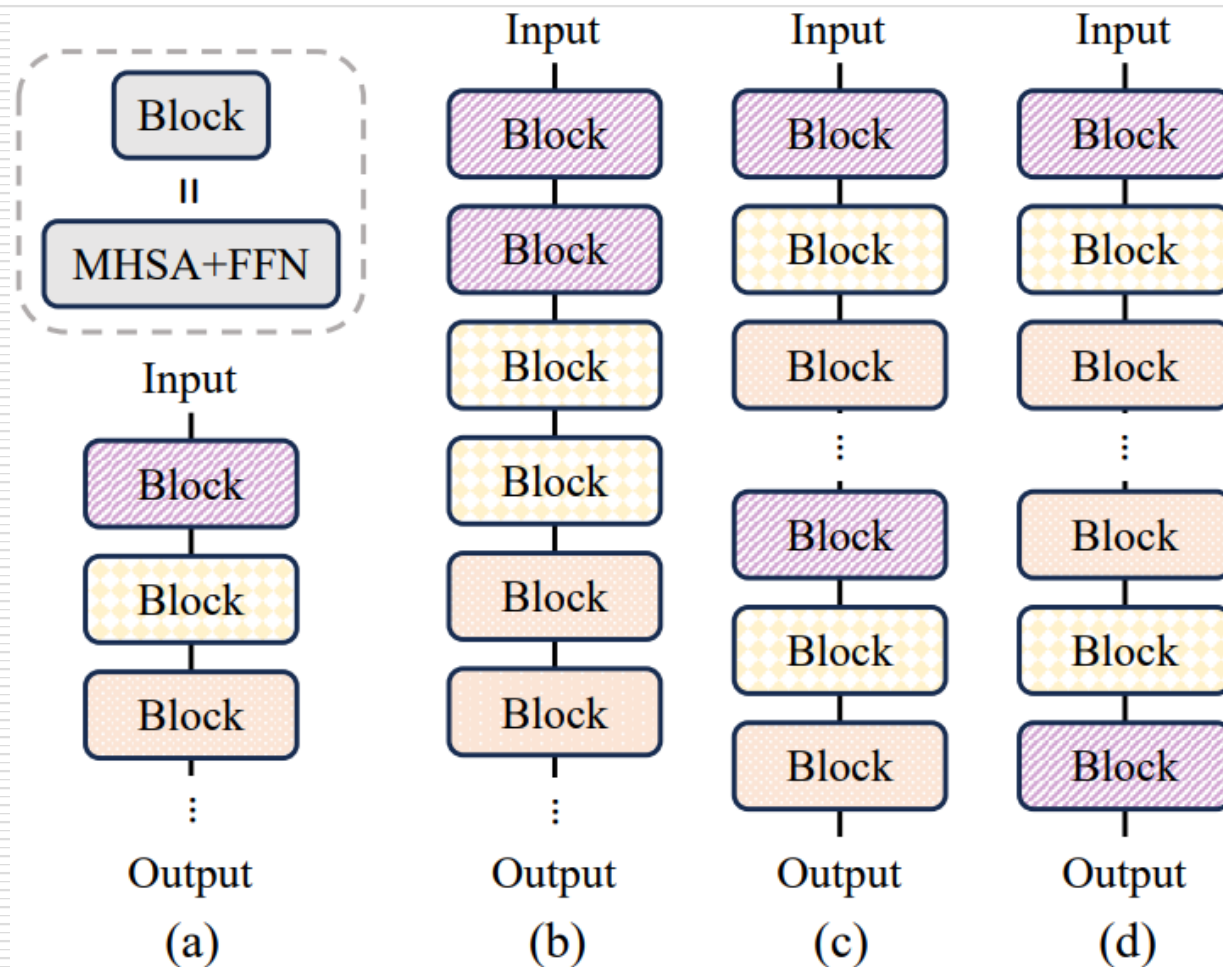
- Weight-sharing by **clustering**
 - Cluster weights and use **centroid** in the indexed array
 - › E.g., k-means
 - Fine-tune the centroid weights with a **summation** of corresponding **gradients**



Weight-Sharing in Language Models (1/3)


- Scaling LMs to achieve better performance
 - Layer-sharing

➤ [MobileLLM](#)




Weight-Sharing in Language Models (2/3)

- Embedding-sharing



Weight-Sharing in Language Models (3/3)

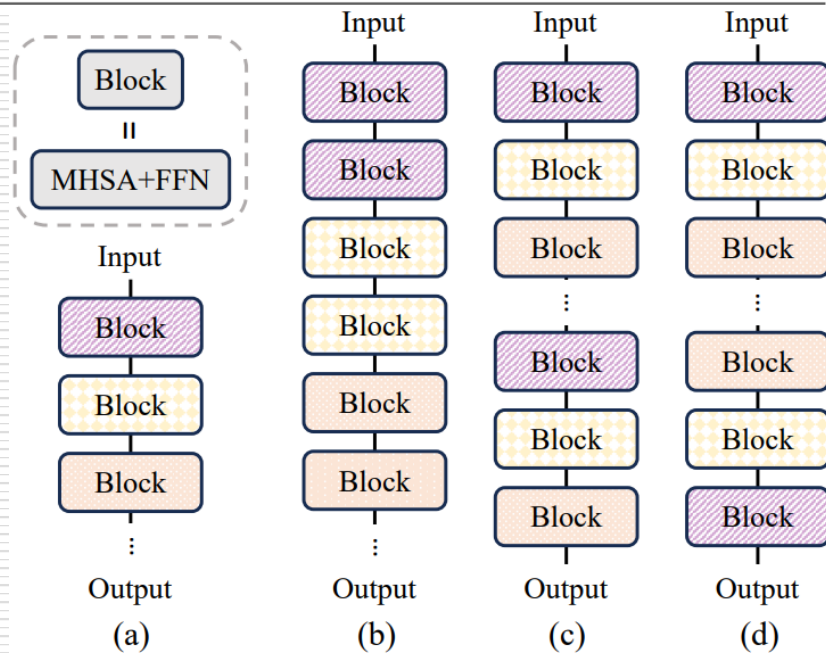
- Embedding-sharing



Experiment Results in MobileLLM

Table 2: Ablation study of layer-sharing strategy on zero-shot common sense reasoning tasks.

Model	Sharing method	ARC-e	ARC-c	BoolQ	PIQA	SIQA	HellaSwag	OBQA	WinoGrande	Avg.
125M	baseline	41.6	25.7	61.1	62.4	43.1	34.4	36.9	51.6	44.6
	Immediate block-wise share	43.9	27.9	61.5	64.3	41.5	35.5	35.1	50.2	45.0
	Repeat-all-over share	43.6	27.1	60.7	63.4	42.6	35.5	36.9	51.7	45.2
	Reverse share	43.8	26.0	58.9	62.9	42.2	35.2	36.8	52.2	44.8
350M	baseline	50.8	30.6	62.3	68.6	43.5	45.1	43.8	52.4	49.6
	Immediate block-wise share	51.5	30.8	59.6	68.2	43.9	47.7	44.7	55.0	50.2
	Repeat-all-over share	53.5	33.0	61.2	69.4	43.2	48.3	42.2	54.6	50.7
	Reverse share	50.7	32.2	61.0	68.8	43.8	47.4	43.1	53.8	50.1



Low-Rank Approximation

Matrix Decomposition

- Any matrix A of rank- r can be decomposed into a long and skinny matrix times a short and long one

The diagram illustrates the matrix decomposition $A = Y Z^T$. Matrix A is a large rectangle with height m and width n . Matrix Y is a tall, narrow rectangle with height m and width r . Matrix Z^T is a short, wide rectangle with height r and width n . The equation $A = Y \times Z^T$ is shown, with the multiplication symbol \times placed between Y and Z^T .

Singular Value Decomposition (SVD)

- A : an $m \times n$ matrix of rank- r
- U : an $m \times m$ orthogonal matrix
- V : an $n \times n$ orthogonal matrix
- S : an $m \times n$ diagonal matrix with nonnegative entries, and with the diagonal entries sorted from high to low

$$A = USV^T$$





Figure 2: The singular value decomposition (SVD). Each singular value in S has an associated left singular vector in U , and right singular vector in V .

Low-Rank Approximation with SVD (1/2)

- Rank-k approximation, $k \leq r$ $A_k = U_k S_k V_k^T$
 - Compute the SVD
 - Set U_k equal to the first k columns of U ($m \times k$)
 - Set S_k equal to the first k rows and columns of S ($k \times k$)
 - Set V_k^T equal to the first k rows of V^T ($k \times n$)




Low-Rank Approximation with SVD (2/2)

- Get $A_k = U_k S_k V_k^T$
- Fuse S_k in U_k or V_k^T
 - $A_k = U'_k V_k^T$ or $U_k V_k^{T'}$
- # of parameters
 - Original
 - › $m \times n$
 - After low-rank approximation
 - › $m \times k + k \times n$

Quantization Concepts


Common Data Types (1/2)

- In neural network
 - Parameters are usually stored in 16-bit or 32-bit precision
 - Single-precision IEEE 754 (FP32)




- Floating point format with lower bits

- › FP16

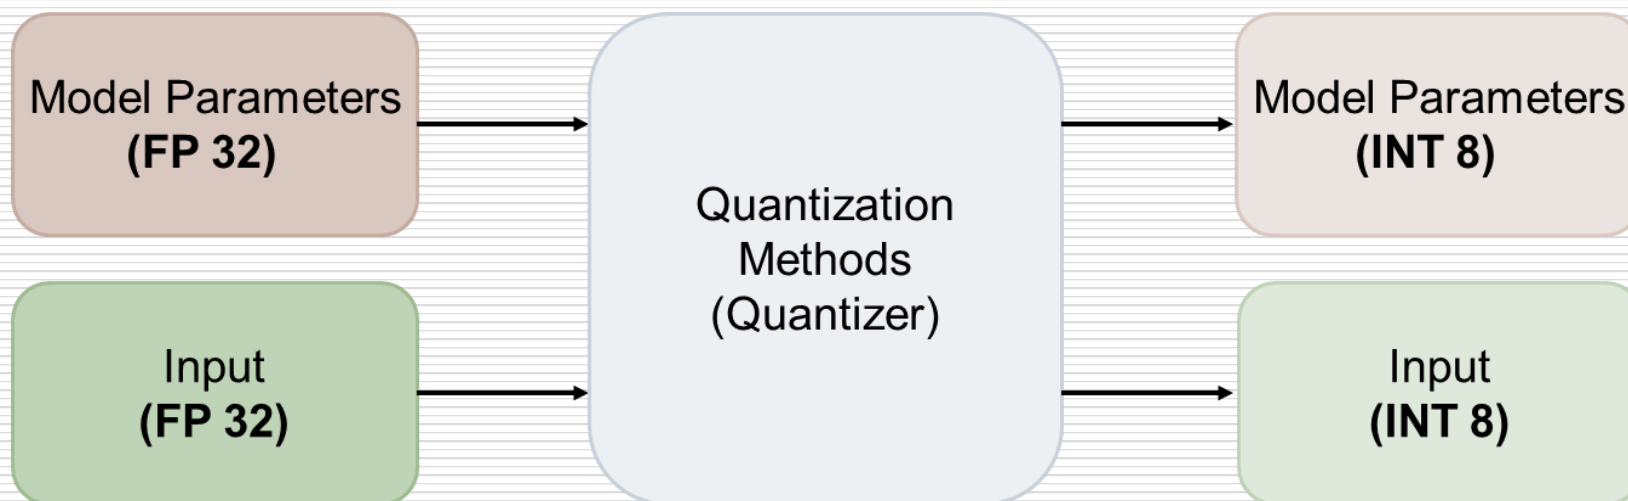


- › FP8



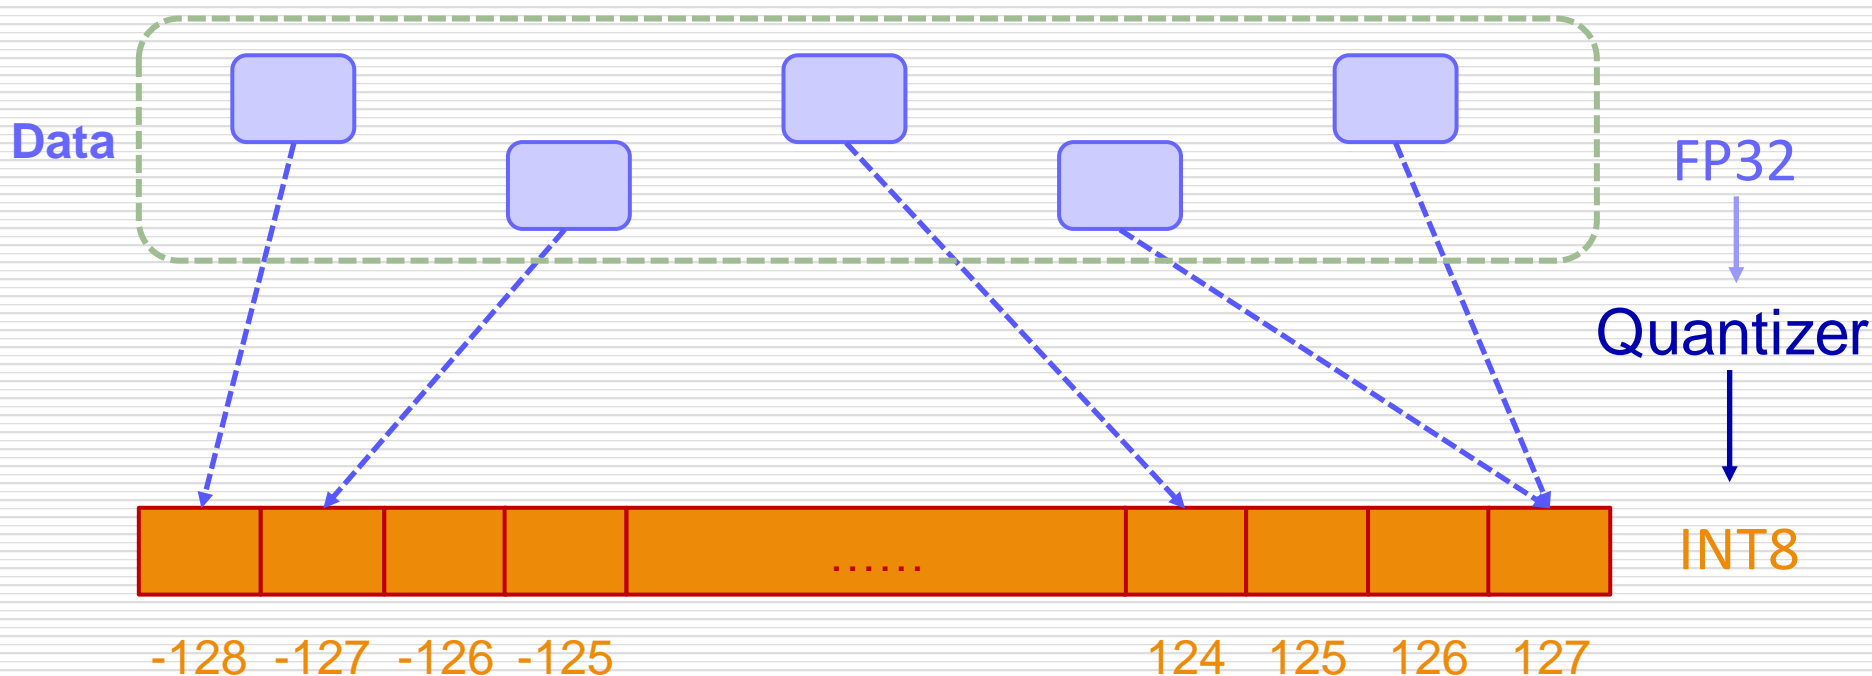
Common Data Types (2/2)

- When quantize higher bits to lower bits
 - Lower computation overhead 😊
 - Reduce memory usage 😊
 - Lower power consumption 😊
 - Quantization error → loss of precision 😞
- FP32 → INT8



Example: FP32 to INT8

- Aim to move values from **FP32** to **INT8**
- Quantizer needs to allocate all FP32 values into 256 states
 - 8 bits can represent 256 states




Uniform Integer Quantization

Basics of N-Bit Uniform Integer Quantization


- Formula

$$X^{INT_N} = \left\lfloor \frac{X^{FP}}{s} \right\rfloor + z, \quad X^{FP} \approx s(X^{INT_N} - z)$$


- Example of 8-bit quantization



Data **before** quantization



Data after **symmetric** quantization



Data after **asymmetric** quantization

Symmetric & Asymmetric Quantization

- Formula

- $X^{INT_N} = \left\lfloor \frac{X^{FP}}{s} \right\rfloor + z$
- $X^{FP} \approx s(X^{INT_N} - z)$ (dequantize)
- $a = -2^{N-1}, b = 2^{N-1} - 1$

s: scale

z: zero point

N: target number of bits

- Symmetric** quantization


- $s_{sym} = \frac{\max(|X^{FP}|)}{b-a}, \quad z_{sym} = 0$

- Asymmetric** quantization

- $s_{asym} = \frac{\max(X^{FP}) - \min(X^{FP})}{b-a}$


- $z_{tmp} = a - \frac{\min(X^{FP})}{s_{asym}}$

- $z_{asym} = \text{clamp}\{z_{tmp}, a, b\}$



Quantization Granularity

Scale Calibration Range



coarse-grained

fine-grained



Accuracy



Efficient



Per-tensor Quantization for Matrix Operation

$$I^{INT} = \left\lfloor \frac{I^{FP}}{s_i} \right\rfloor + z_i, \quad \overline{I^{FP}} = s_i(I^{INT} - z_i)$$

$$W^{INT} = \left\lfloor \frac{W^{FP}}{s_w} \right\rfloor + z_w, \quad \overline{W^{FP}} = s_w(W^{INT} - z_w)$$

$$B^{INT} = \left\lfloor \frac{B^{FP}}{s_i \cdot s_w} \right\rfloor, \quad \overline{B^{FP}} = s_i s_w \cdot B^{INT}$$

$$O = I \cdot W + B \approx \overline{O^{FP}} = \overline{I^{FP}} \cdot \overline{W^{FP}} + \overline{B^{FP}}, \quad O^{INT} = \left\lfloor \frac{O^{FP}}{s_o} \right\rfloor + z_o$$

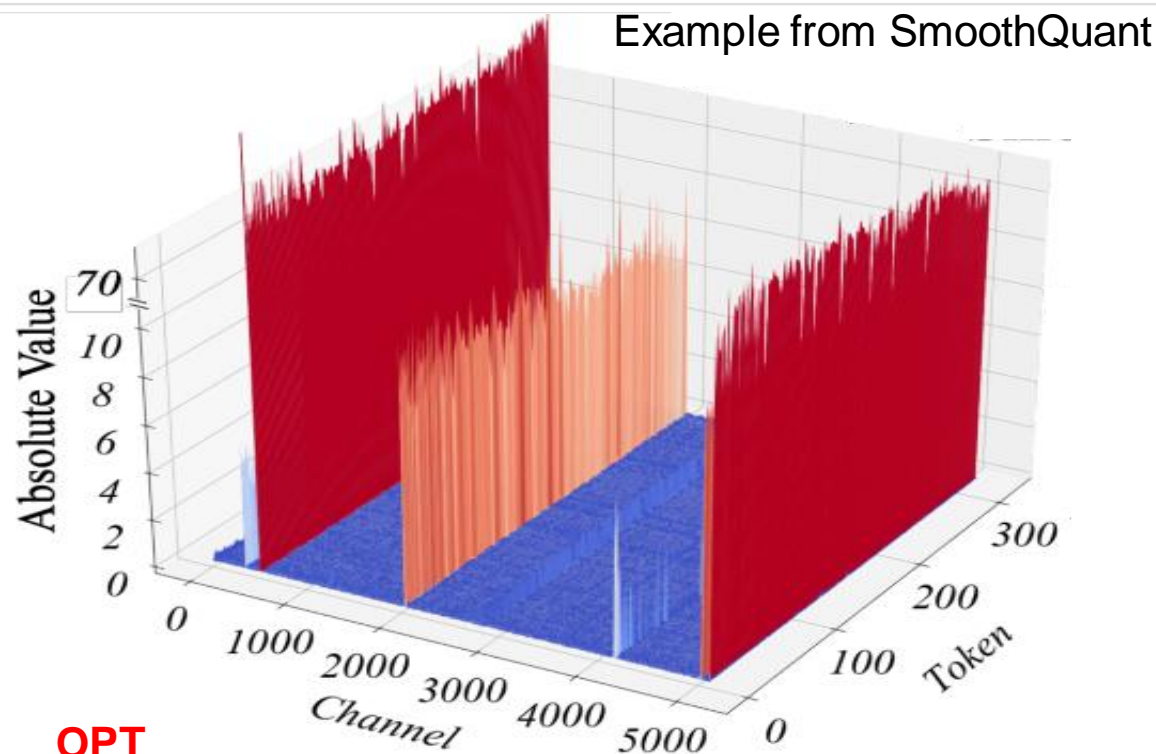
$$O^{INT} = \left\lfloor \frac{\sum s_i(I^{INT} - z_i)s_w(W^{INT} - z_w) + s_i s_w \cdot B^{INT}}{s_o} \right\rfloor + z_o = \left\lfloor \frac{s_i s_w}{s_o} \left[\sum (I^{INT} - z_i)(W^{INT} - z_w) + B^{INT} \right] \right\rfloor + z_o$$

$$\frac{s_i s_w}{s_o} \approx M^{INT} \cdot 2^{shift}$$

Survey of LLM Quantization

Quantization Difficulty of LLMs (1/2)

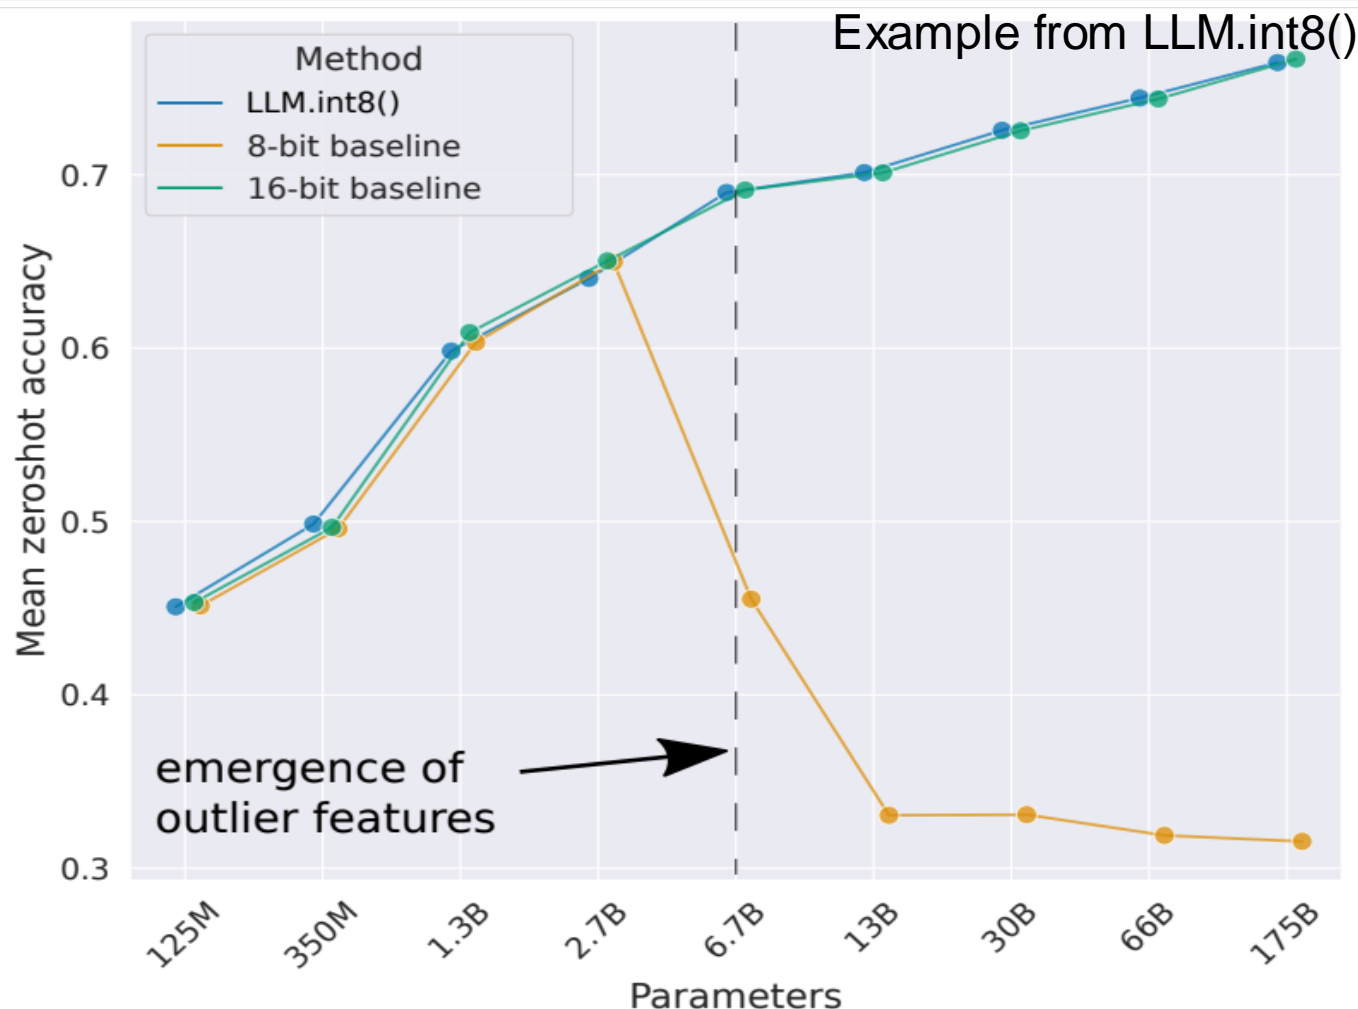
- Outliers in OPT (SmoothQuant)



OPT					
Model	#L	#H	d_{model}	LR	Batch
13B	40	40	5120	$1.0e-4$	4M


Quantization Difficulty of LLMs (2/2)

- Apply the vanilla 8-bit quantization on LLM



LLM.int8(): The Fact of Outliers (1/3)

- These outlier features are **highly systematic** after **the emergence of outlier features** occurs
 - EX : 6.7B transformer with a sequence length of 2048
 - › About 150k outlier features per sequence for the entire transformer, but these features are concentrated in **only 7 different hidden dimensions**



LLM.int8(): The Fact of Outliers (2/3)

- The effect of outlier features
 - Set all outlier features (at most **7** hidden dimensions) of a layer to zero
 - › 為了不要讓error累積，一次只做一層，其他層保持原樣
 - › Accuracy會掉20~40%
 - › Perplexity increases by 600~1000%
 - Set 7 random feature dimensions of a layer to zero
 - › 為了不要讓error累積，一次只做一層，其他層保持原樣
 - › Accuracy只掉0.02~0.3%
 - › Perplexity increases by 0.1%


LLM.int8(): The Fact of Outliers (3/3)

- The effect of outlier features
 - These outliers are critical for transformer performance
 - Quantization precision for these outlier features is paramount as even tiny errors greatly impact model performance

LLM.int8(): Methodology

- Structured outliers


LLM.int8()



Regular values
Outliers

8-bit Vector-wise Quantization

(1) Find vector-wise constants: C_W & C_X



(2) Quantize

$$X_{F16} * (127/C_X) = X_{I8}$$

$$W_{F16} * (127/C_W) = W_{I8}$$

(3) Int8 Matmul


$$X_{I8} W_{I8} = Out_{I32}$$

(4) Dequantize

$$\frac{Out_{I32} * (C_X \otimes C_W)}{127 * 127} = Out_{F16}$$

16-bit Decomposition

(1) Decompose outliers



(2) FP16 Matmul

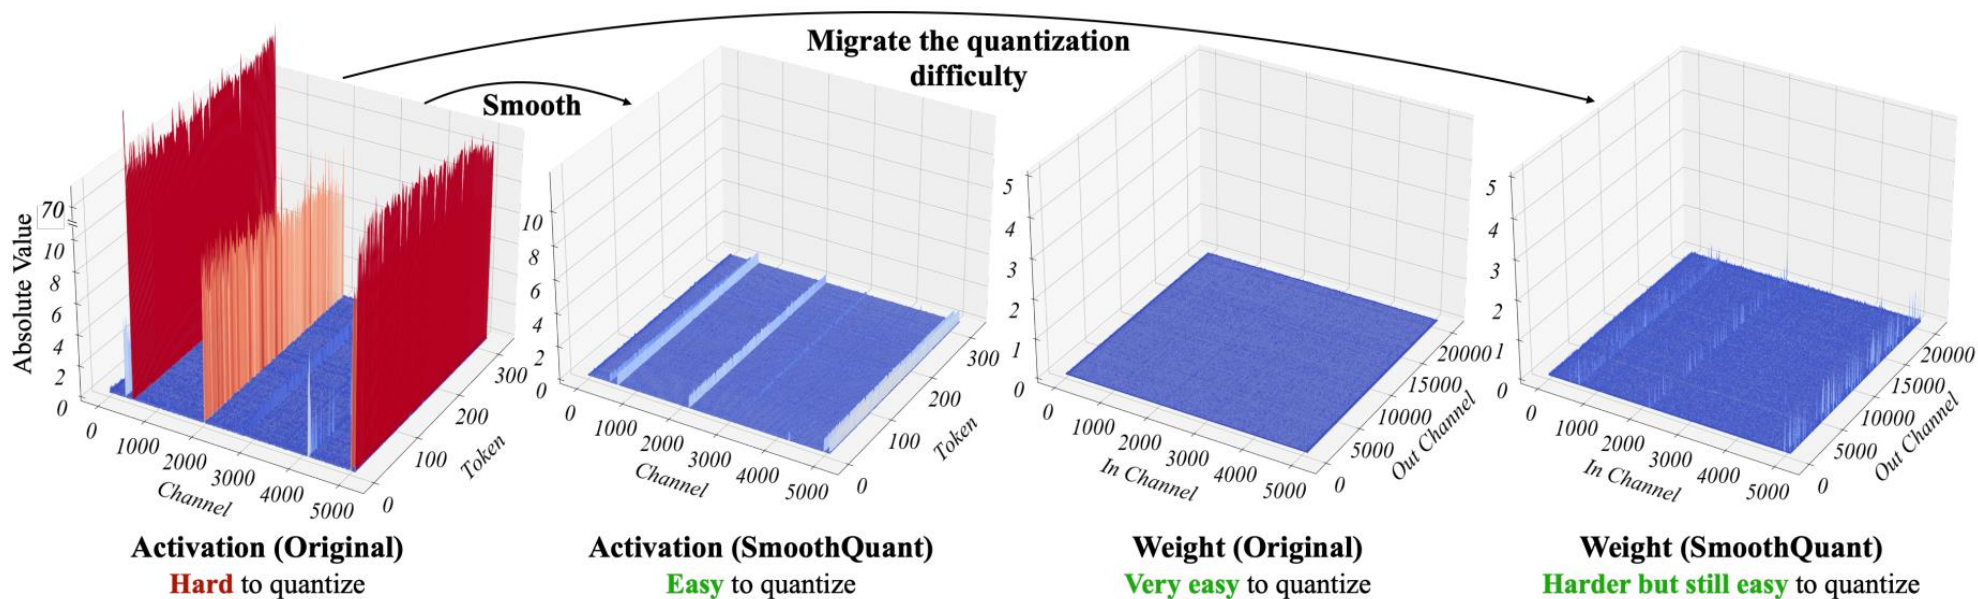
$$X_{F16} W_{F16} = Out_{F16}$$



Out_{FP16}

SmoothQuant (1/3)


- Idea
 - Migrate the quantization difficulty from activations to weights



SmoothQuant (2/3)

- Channel-wise smoothing

$$\mathbf{O} = (\mathbf{I} \cdot \text{diag}(\lambda)^{-1}) \cdot (\text{diag}(\lambda) \cdot \mathbf{W}) = \hat{\mathbf{I}}\hat{\mathbf{W}}$$

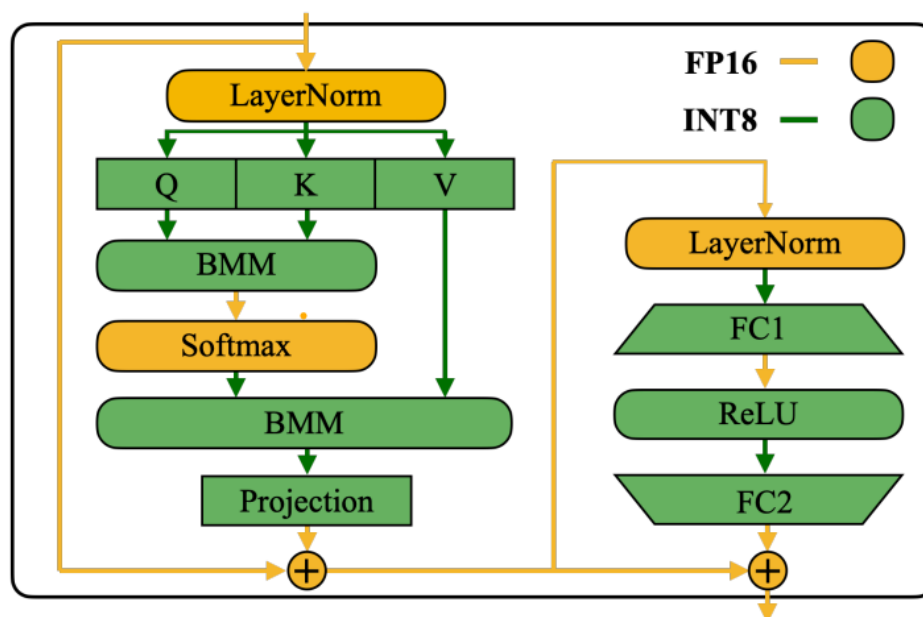


SmoothQuant (3/3)

- Propose **smoothing factor** $\lambda \in \mathbb{R}^{C_i}$
 - Splits some quantization difficulties to the weights

$$\lambda_\rho = \frac{\maxAbs(I_\rho)^\alpha}{\maxAbs(W_\rho)^{1-\alpha}}, \quad \rho \in 1, 2, \dots, C_{embd}$$

- α control how much difficulty is migrated



FPTQ

- Logarithmic Activation Equalization

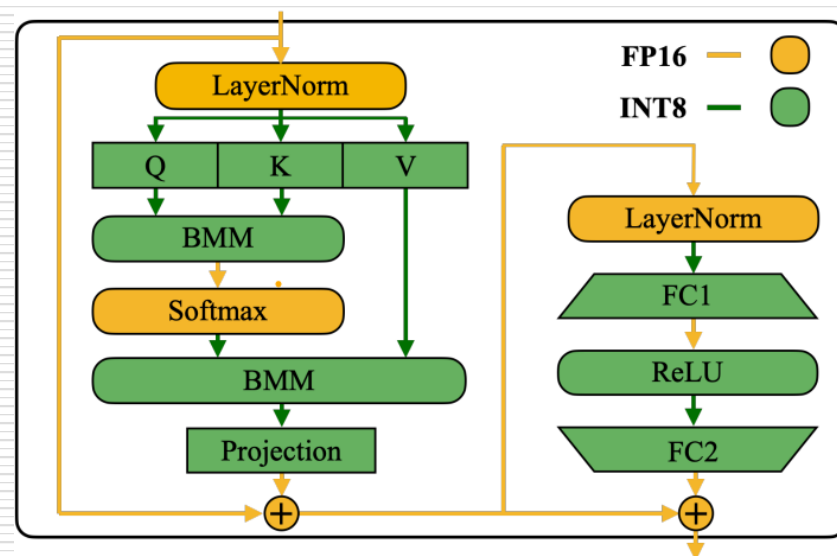

- Weight-independent

$$\lambda_\rho = \frac{\maxAbs(\mathbf{I}_\rho)}{\log_2(2 + \maxAbs(\mathbf{I}_\rho))}, \quad \rho \in 1, 2, \dots, C_{embd}$$

- If outlier exceed a preset threshold


- Dynamic quantization

$$\mathbf{O} = (\mathbf{I} \cdot \text{diag}(\lambda)^{-1}) \cdot (\text{diag}(\lambda) \cdot \mathbf{W}) = \hat{\mathbf{I}}\hat{\mathbf{W}}$$



RPTQ (1/2)

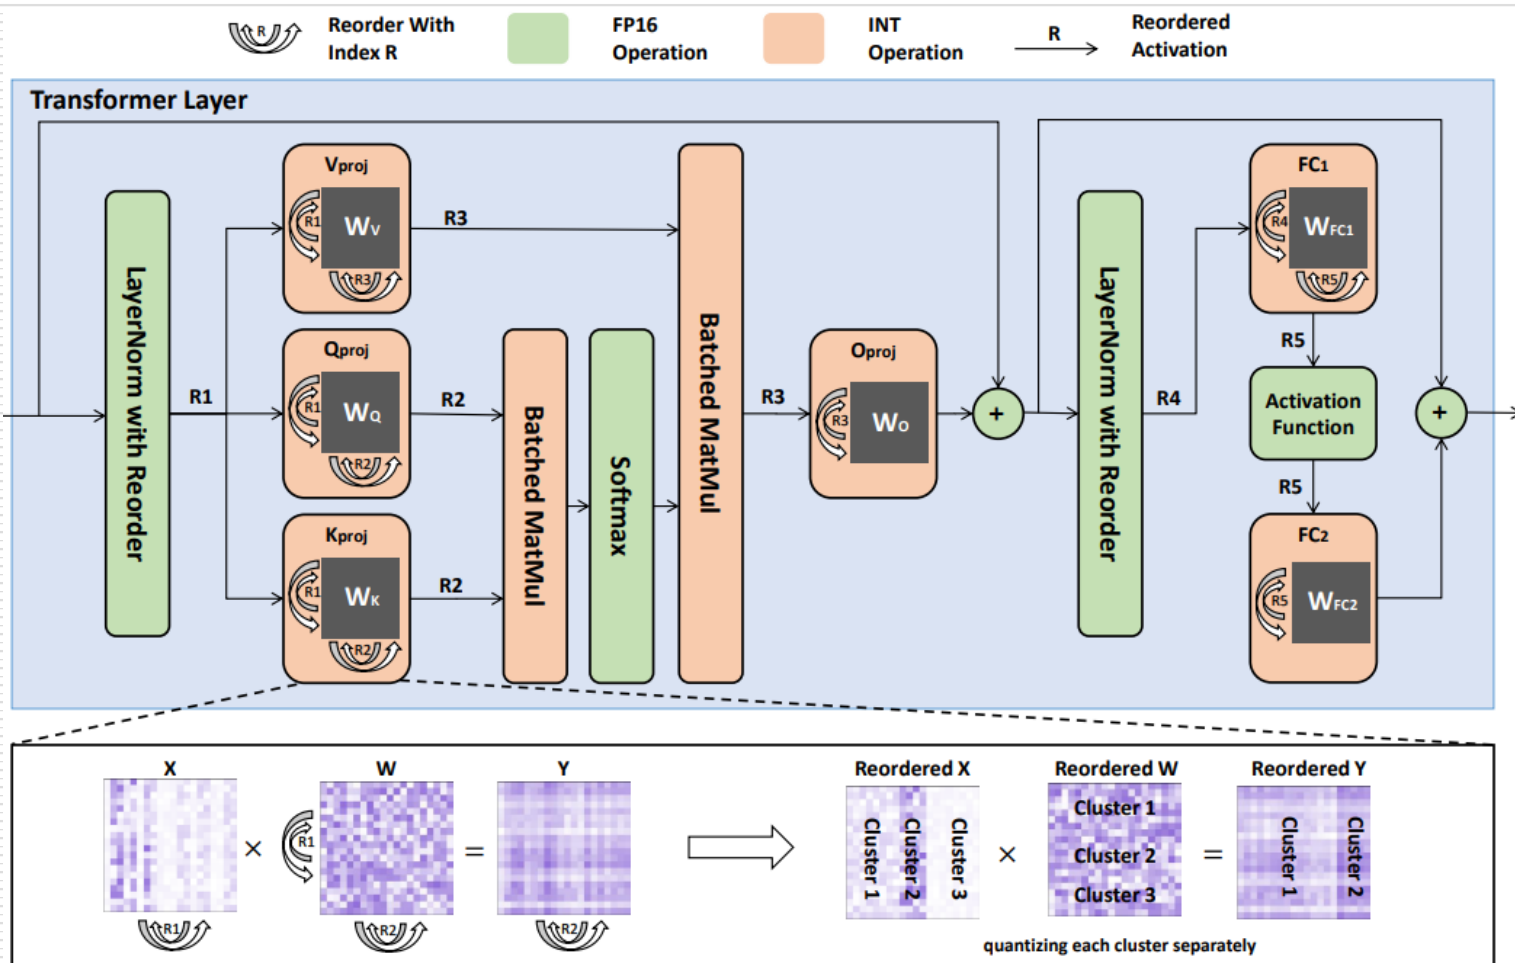
- Per-group quantization
 - Reorder activation & weight by clustering algorithm



- Flow
 - Representative vector
 - › Each activation channel's (min, max) values after calibration
 - Apply k-means, then get the result R
 - Activations and their corresponding weights are reordered by R to maintain computational equivalence
 - Different scales are calculated for different clusters

RPTQ (2/2)

- Per-group quantization
 - Reorder activation & weight by clustering algorithm



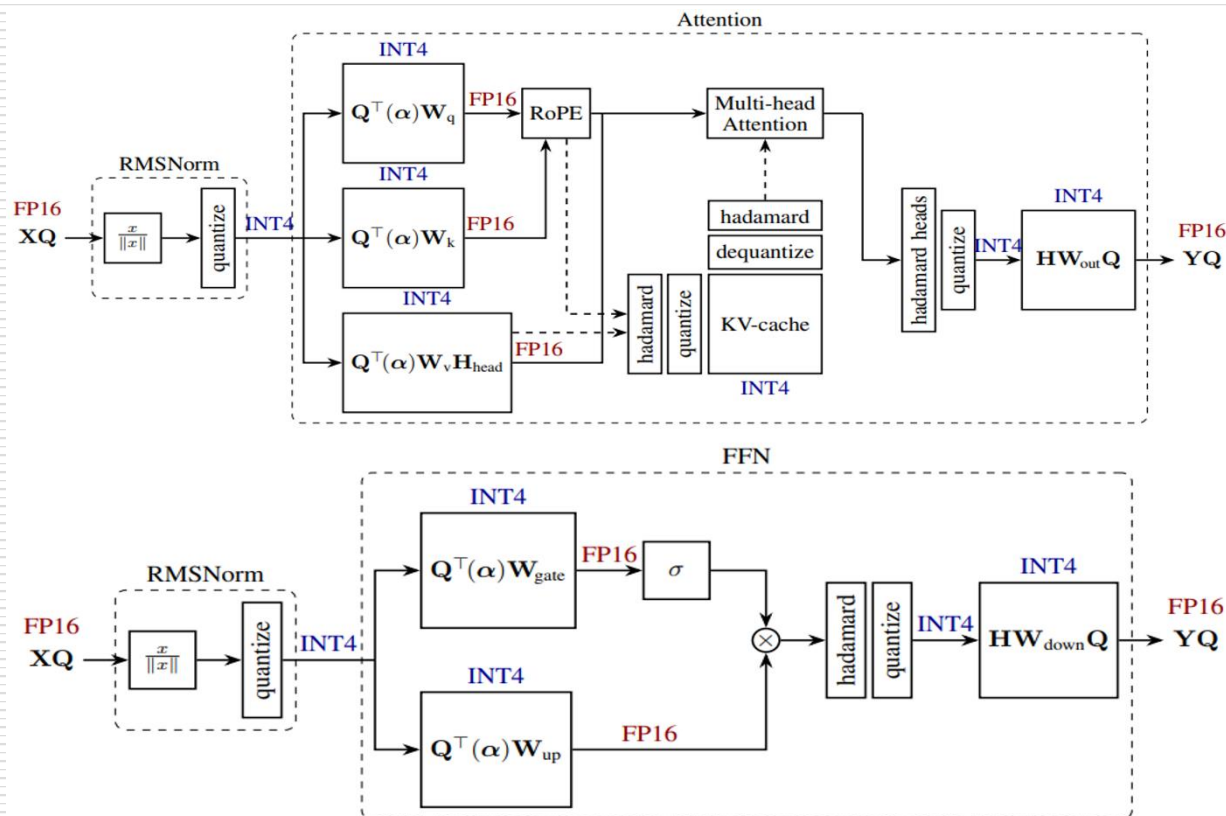
QuaRot (1/2)

- Rotation-based



QuaRot (2/2)

- Retains floating-point operations for
 - Layer normalization, residual link additions
 - Rotation of input activations
 - Multi-head Attention
 - Gating in the FFN



PFPQuant:

A Pure Fixed-Point Quantization Framework for Edge Devices without Floating-Point Units

LLM Quantization w/o FPU

- FPU
 - Floating-point units
- w/o FPU
 - Means all computation should be
 - › Integer operation & static shifts
- Difficulty
 - Most of uniform integer quantization is mixed precision
 - › **Layer normalization** (LayerNorm 、 RMSNorm)
 - » All
 - › Residual-link addition
 - » All
 - › Outlier-related computation
 - » LLM.int8()



Main Components of PFPQuant

- Generalized Outlier Mitigation
 - GOM
- Outlier-Reduced Integer LayerNorm
 - ORI-LN
- Residual-Link with Space-Accuracy Tradeoff
 - RL-SAT

GOM: Outlier Mitigation (1/2)

- Equivalence

$$O = I \cdot W = \underbrace{\{I \cdot \text{diag}(\lambda)^{-1}\}}_{I_{\text{new}}} \cdot \underbrace{\{\text{diag}(\lambda) \cdot W\}}_{W_{\text{new}}}$$

- **E**nhanced **L**ogarithmic **A**ctivation **E**qualization
 - E-LAE
 - **LAE** proposed by FPTQ lacks the flexibility to adjust various quantization nodes
 - We introduced an **α** exponent term to the original formula

$$\lambda_{\rho} = \left[\frac{\maxAbs(I_{\rho})}{\log_2(2 + \maxAbs(I_{\rho}))} \right]^{\alpha}, \quad \rho \in 1, 2, \dots, C_{\text{embd}}$$

GOM: Outlier Mitigation (2/2)

- **M**edian-**C**onvergent **O**utlier **M**itigation
 - MCOM
 - Goal
 - › Static per-tensor quantization for input activations
 - » Extremely hardware-friendly
 - Idea
 - › Based on **E-LAE**
 - › Concentrate the elements of the same tensor around the median

$$V_{maxAbs} = \{maxAbs(I_1), \quad maxAbs(I_2), \quad \dots, \quad maxAbs(I_{C_{embd}})\}$$

$$\lambda_\rho = \left[\frac{\beta + maxAbs(I_\rho)}{\log_2(2 + median(V_{maxAbs}))} \right]^\alpha, \quad \rho \in 1, 2, \dots, C_{embd}$$

GOM: Combined with Data-Shifting

- Data-shifting vector

$$\delta_{\rho} = \frac{-1}{2} [\max(I_{\rho}) + \min(I_{\rho})], \quad \rho \in 1, 2, \dots, C_{embd}$$

- Final formula of **GOM**

$$O = \boxed{\{(I + \delta) \cdot \text{diag}(\lambda)^{-1}\}} \cdot \boxed{\{\text{diag}(\lambda) \cdot W\}} + \boxed{\{B - \delta \cdot W\}}$$

I_{easyQ} ← (points to the green box)

← (points to the blue box) *W_{new}*

← (points to the red box) *B_{new}*


ORI-LN: Outlier Score (1/2)

- Previous quantization research has not performed LayerNorm under quantized precision
- To elucidate the difficulty of LayerNorm quantization
 - We define an outlier score ***OS***
 - › After calibration, V_{maxAbs} is a static vector

$$V_{maxAbs} = \{maxAbs(I_1), \quad maxAbs(I_2), \quad \dots, \quad maxAbs(I_{C_{embd}})\}$$

$$OS = \frac{\mathbf{max}(V_{maxAbs})}{\mathbf{median}(V_{maxAbs})}$$

ORI-LN: Outlier Score (2/2)



ORI-LN: Vanilla Integer LayerNorm (1/2)

- Static per-tensor quantization for input activations
 - For simplicity, we use the RMSNorm formula, a variant of LayerNorm, to illustrate our approach

$$O_{\rho} = \frac{I_{\rho} \cdot \gamma_{\rho}}{\sqrt{\frac{1}{C_{embd}} \cdot \sum_{m=1}^{C_{embd}} (I_m)^2}}, \quad \rho \in 1, 2, \dots, C_{embd}$$

$$I_{\rho}^{INT} = \left\lfloor \frac{I_{\rho}}{s_i} \right\rfloor + z_i, \quad I_{\rho} \approx s_i (I_{\rho}^{INT} - z_i) \quad \gamma_{\rho} \approx \gamma_{\rho}^{INT} \cdot s_{\gamma_{\rho}}, \quad s_{\gamma_{\rho}} = 2^{shift_{\gamma_{\rho}}}$$

$$O_{\rho} \approx \frac{s_i (I_{\rho}^{INT} - z_i) \cdot s_{\gamma_{\rho}} \cdot \gamma_{\rho}^{INT}}{\sqrt{\frac{1}{C_{embd}} \cdot \sum_{m=1}^{C_{embd}} (s_i (I_m^{INT} - z_i))^2}} = \frac{\cancel{s_i}}{\sqrt{\cancel{s_i^2}}} \cdot (s_{\gamma_{\rho}} \cdot \sqrt{C_{embd}}) \cdot \frac{(I_{\rho}^{INT} - z_i) \cdot \gamma_{\rho}^{INT}}{\sqrt{\sum_{m=1}^{C_{embd}} ((I_m^{INT} - z_i))^2}}$$

$$s_{\gamma_{\rho}} \cdot \sqrt{C_{embd}} \approx M_{\rho}^{INT} \cdot 2^{shift_{M_{\rho}}}$$

ORI-LN: Vanilla Integer LayerNorm (2/2)

- SQ represents the outlier mitigation formula proposed by SmoothQuant
- Using the LLaMA-2 7B chat model
 - The average accuracy of each method on LAMBADA, HellaSwag, PIQA, WinoGrande, and ARCe
 - The perplexity on Wikitext-2

Method	Average↑	Wikitext↓
FP	0.7226	12.28
SQ w VILN	0.3106	51036.37
E-LAE w VILN	0.3092	53769.12

ORI-LN

- Introducing outlier mitigation into LayerNorm

$$O_{\rho} = \frac{I_{\rho} \cdot \gamma_{\rho}}{\sqrt{\frac{1}{C_{embd}} \cdot \sum_{m=1}^{C_{embd}} (I_m)^2}} = \frac{\sqrt{C_{embd}}}{\sqrt{\sum_{m=1}^{C_{embd}} \left(\frac{I_m}{\lambda_m} \cdot \lambda_m \right)^2}} \cdot \left(\frac{I_{\rho}}{\lambda_{\rho}} \right) \cdot (\lambda_{\rho} \cdot \gamma_{\rho}), \quad \rho \in 1, 2, \dots, C_{embd}$$


$$O_{\rho} = \frac{\sqrt{C_{embd}}}{\sqrt{\sum_{m=1}^{C_{embd}} \left(\frac{I_m}{\lambda_m} \right)^2 \cdot (\lambda_m)^2}} \cdot \left(\frac{I_{\rho}}{\lambda_{\rho}} \right) \cdot (\lambda_{\rho} \cdot \gamma_{\rho}) \quad I'_{\rho} = \left(\frac{I_{\rho}}{\lambda_{\rho}} \right), \quad \gamma'_{\rho} = (\lambda_{\rho} \cdot \gamma_{\rho})$$

$$O_{\rho} = \frac{I'_{\rho} \cdot \gamma'_{\rho} \cdot \sqrt{C_{embd}}}{\sqrt{\sum_{m=1}^{C_{embd}} (I'_m)^2 \cdot (\lambda_m)^2}} \approx \frac{s_{i'} (I'^{INT}_{\rho} - z_{i'}) \cdot s_{\gamma'_{\rho}} \cdot \gamma'^{INT}_{\rho} \cdot \sqrt{C_{embd}}}{\sqrt{\sum_{m=1}^{C_{embd}} (s_{i'} (I'^{INT}_m - z_{i'}))^2 \cdot (\lambda_m)^2}} = (s_{\gamma'_{\rho}} \cdot \sqrt{C_{embd}}) \cdot \frac{(I'^{INT}_{\rho} - z_{i'}) \cdot \gamma'^{INT}_{\rho}}{\sqrt{\sum_{m=1}^{C_{embd}} (I'^{INT}_m - z_{i'})^2 \cdot (\lambda_m)^2}}$$

$$s_{\gamma'_{\rho}} \cdot \sqrt{C_{embd}} \approx M_{\rho}^{INT} \cdot 2^{shift_{M_{\rho}}}, (\lambda_m)^2 \approx M_{\lambda_m}^{INT} \cdot 2^{shift_{M_{\lambda_m}}}$$

RL-SAT: Recall (1/2)

- Since the residual-link originates from the input of the previous LayerNorm, it inherits its quantization difficulty



RL-SAT: Recall (2/2)

- Per-tensor quantization for matrix operation

$$I^{INT} = \left\lceil \frac{I^{FP}}{s_i} \right\rceil + z_i, \quad \overline{I^{FP}} = s_i(I^{INT} - z_i)$$

$$W^{INT} = \left\lceil \frac{W^{FP}}{s_w} \right\rceil + z_w, \quad \overline{W^{FP}} = s_w(W^{INT} - z_w)$$

$$B^{INT} = \left\lceil \frac{B^{FP}}{s_i \cdot s_w} \right\rceil, \quad \overline{B^{FP}} = s_i s_w \cdot B^{INT}$$

$$O = I \cdot W + B \approx \overline{O^{FP}} = \overline{I^{FP}} \cdot \overline{W^{FP}} + \overline{B^{FP}}, \quad O^{INT} = \left\lceil \frac{\overline{O^{FP}}}{s_o} \right\rceil + z_o$$

$$O^{INT} = \left\lceil \frac{\sum s_i(I^{INT} - z_i)s_w(W^{INT} - z_w) + s_i s_w \cdot B^{INT}}{s_o} \right\rceil + z_o = \left\lceil \frac{s_i s_w}{s_o} \left[\sum (I^{INT} - z_i)(W^{INT} - z_w) + B^{INT} \right] \right\rceil + z_o$$

$$\frac{s_i s_w}{s_o} \approx M^{INT} \cdot 2^{shift}$$

RL-SAT: Vanilla Residual Quantization

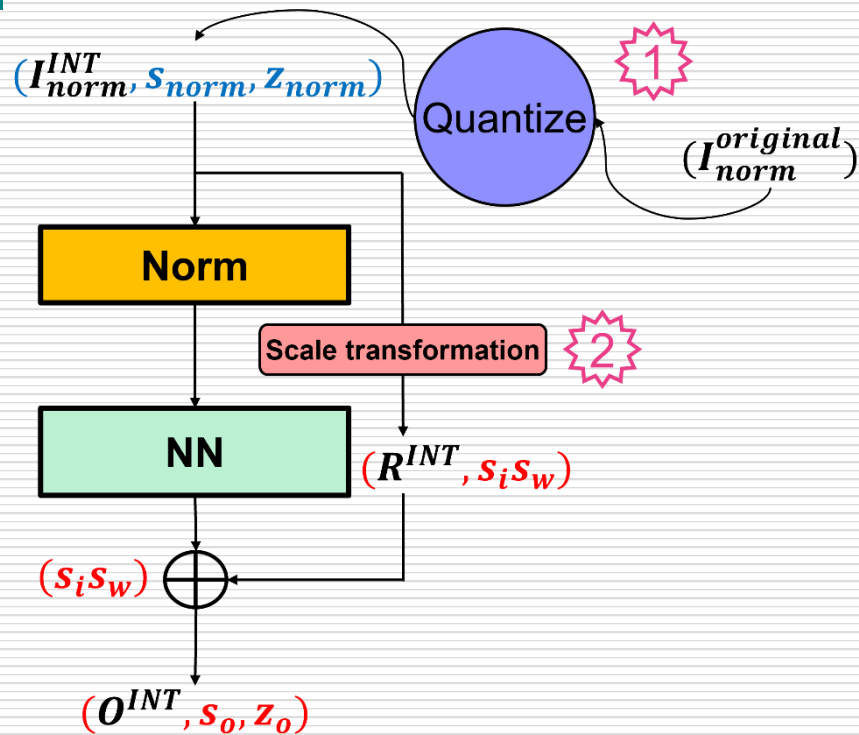
$$I_{norm}^{INT} = \left\lfloor \frac{I_{norm}^{original}}{s_{norm}} \right\rfloor + z_{norm}$$

$$\overline{R_0^{FP}} = \overline{I_{norm}^{FP}} = s_{norm} \cdot (I_{norm}^{INT} - z_{norm}) \approx I_{norm}^{original}$$

$$R^{INT} = \left\lfloor \frac{\overline{R_0^{FP}}}{s_i s_w} \right\rfloor = \left\lfloor \frac{s_{norm}}{s_i s_w} \cdot (I_{norm}^{INT} - z_{norm}) \right\rfloor$$

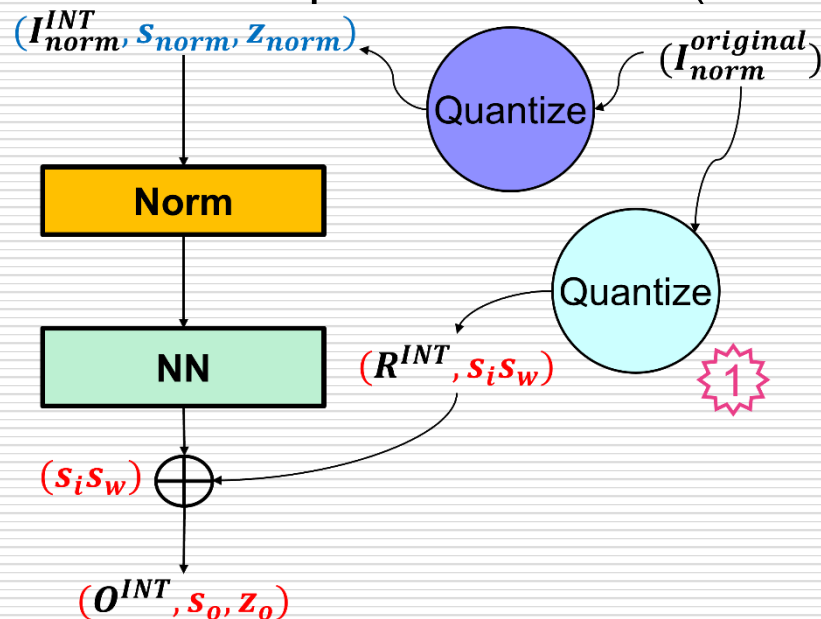
$$\overline{R^{FP}} = s_i s_w \cdot \left\lfloor \frac{s_{norm}}{s_i s_w} \cdot (I_{norm}^{INT} - z_{norm}) \right\rfloor = s_i s_w \cdot R^{INT}$$

$$\frac{s_{norm}}{s_i s_w} \approx M_{scTransform}^{INT} \cdot 2^{shift_{scTransform}} = M_{scTransform}$$



RL-SAT


- Directly quantizes $I_{norm}^{original}$ to a scale that allows correct accumulation
- While this approach uses slightly more memory space, it reduces one instance of quantization error (due to rounding)
- Since the additional memory consumption is negligible compared to the weights of LLMs, PFPQuant adopts this method (RL-SAT)



$$R^{INT} = \left\lfloor \frac{I_{norm}^{original}}{s_i \cdot s_w} \right\rfloor, \quad \overline{R^{FP}} = s_i s_w \cdot R^{INT} \approx R^{FP} = I_{norm}^{original}$$

Experiment Results: Effectiveness of GOM

- GOM
 - Generalized Outlier Mitigation



Experiment Result: Quantized LLaMA 2

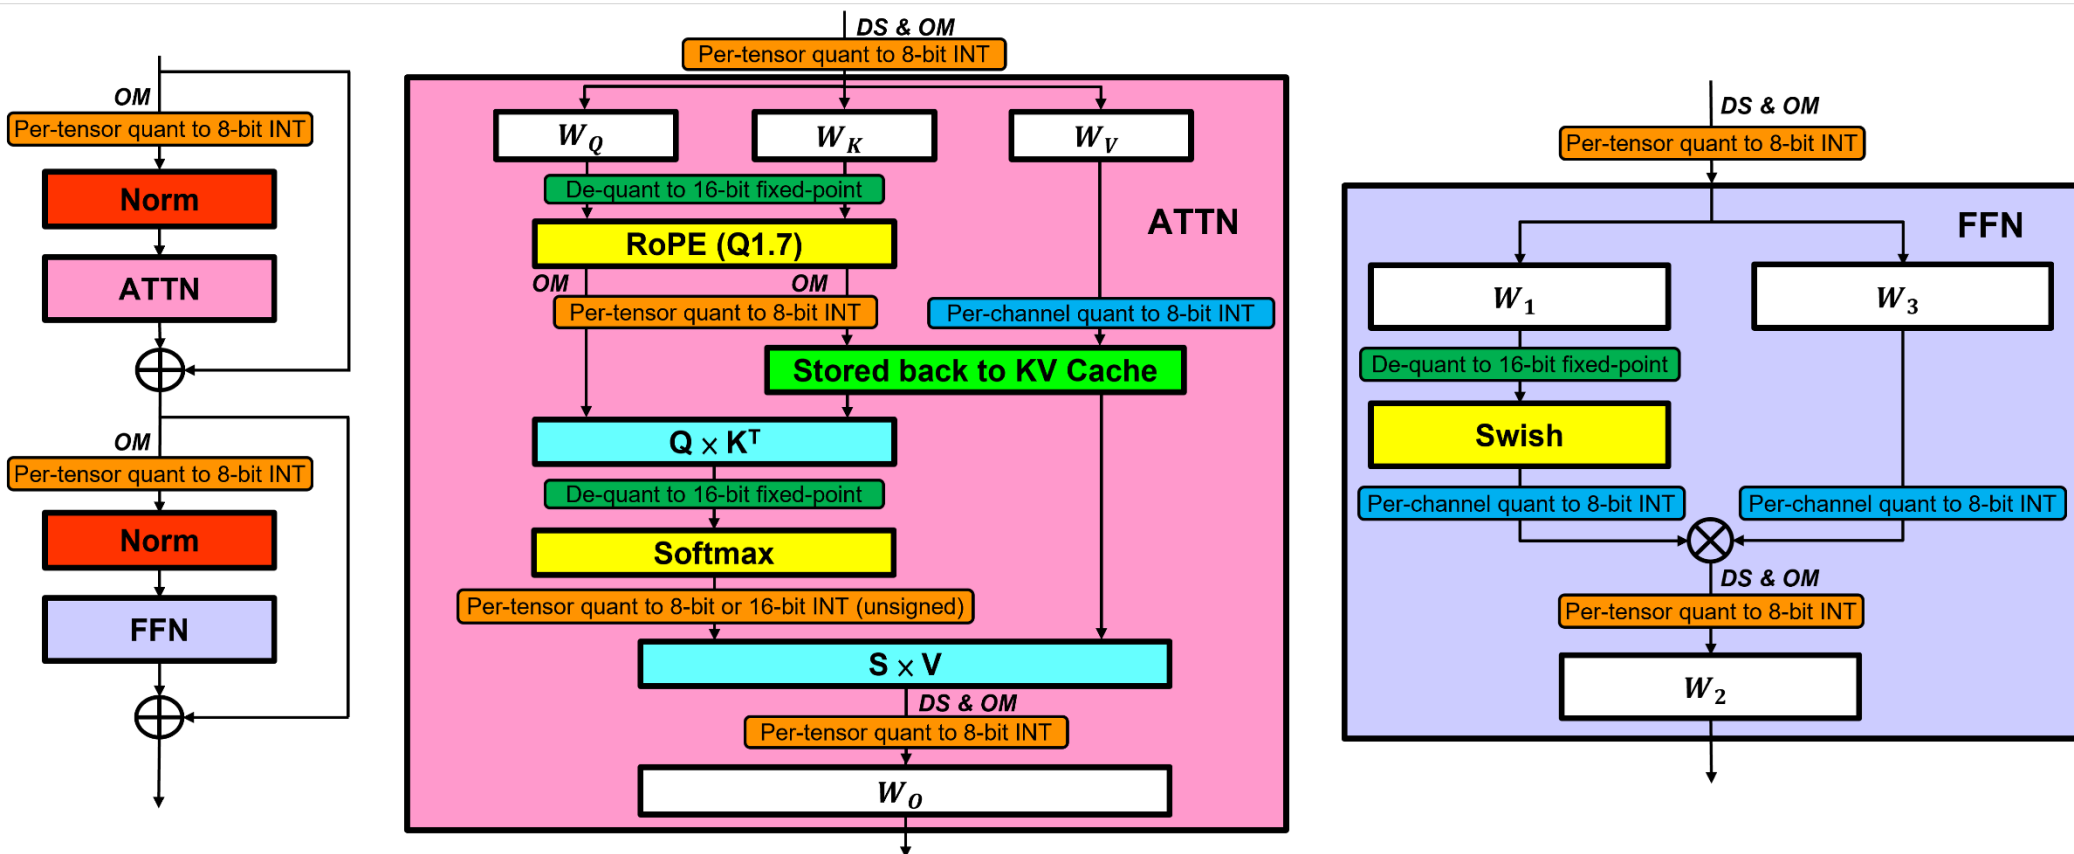
- We tested PFPQuant on various common tasks using the LLaMA-2 7B chat model

Method	LAMBDA	HellaSwag	PIQA	WinoGrande	ARCe	Average↑	Wikitext↓
FP	0.6850	0.7530	0.7666	0.6646	0.7441	0.7226	12.28
SQ w R8-S8	0.6238	0.7266	0.7568	0.6464	0.7033	0.6914	41.52
SQ w R16-S8	0.6286	0.7269	0.7579	0.6409	0.7050	0.6919	41.46
SQ w R8-S16	0.6173	0.7182	0.7546	0.6346	0.7066	0.6863	22.55
SQ w R16-S16	0.6282	0.7220	0.7541	0.6472	0.7104	0.6924	22.39
E-LAE w R8-S8	0.6255	0.7182	0.7541	0.6251	0.7037	0.6853	36.08
E-LAE w R16-S8	0.6271	0.7192	0.7530	0.6172	0.7066	0.6846	36.15
E-LAE w R8-S16	0.6112	0.7111	0.7552	0.6227	0.7075	0.6815	19.34
E-LAE w R16-S16	0.6151	0.7125	0.7519	0.6283	0.7079	0.6831	19.25
MCOM w R8-S8	0.6275	0.7380	0.7519	0.6527	0.7189	0.6978	27.86
MCOM w R16-S8	0.6295	0.7378	0.7486	0.6701	0.7180	0.7008	26.73
MCOM w R8-S16	0.6327	0.7383	0.7552	0.6559	0.7243	0.7013	14.80
MCOM w R16-S16	0.6339	0.7410	0.7579	0.6535	0.7218	0.7016	14.77

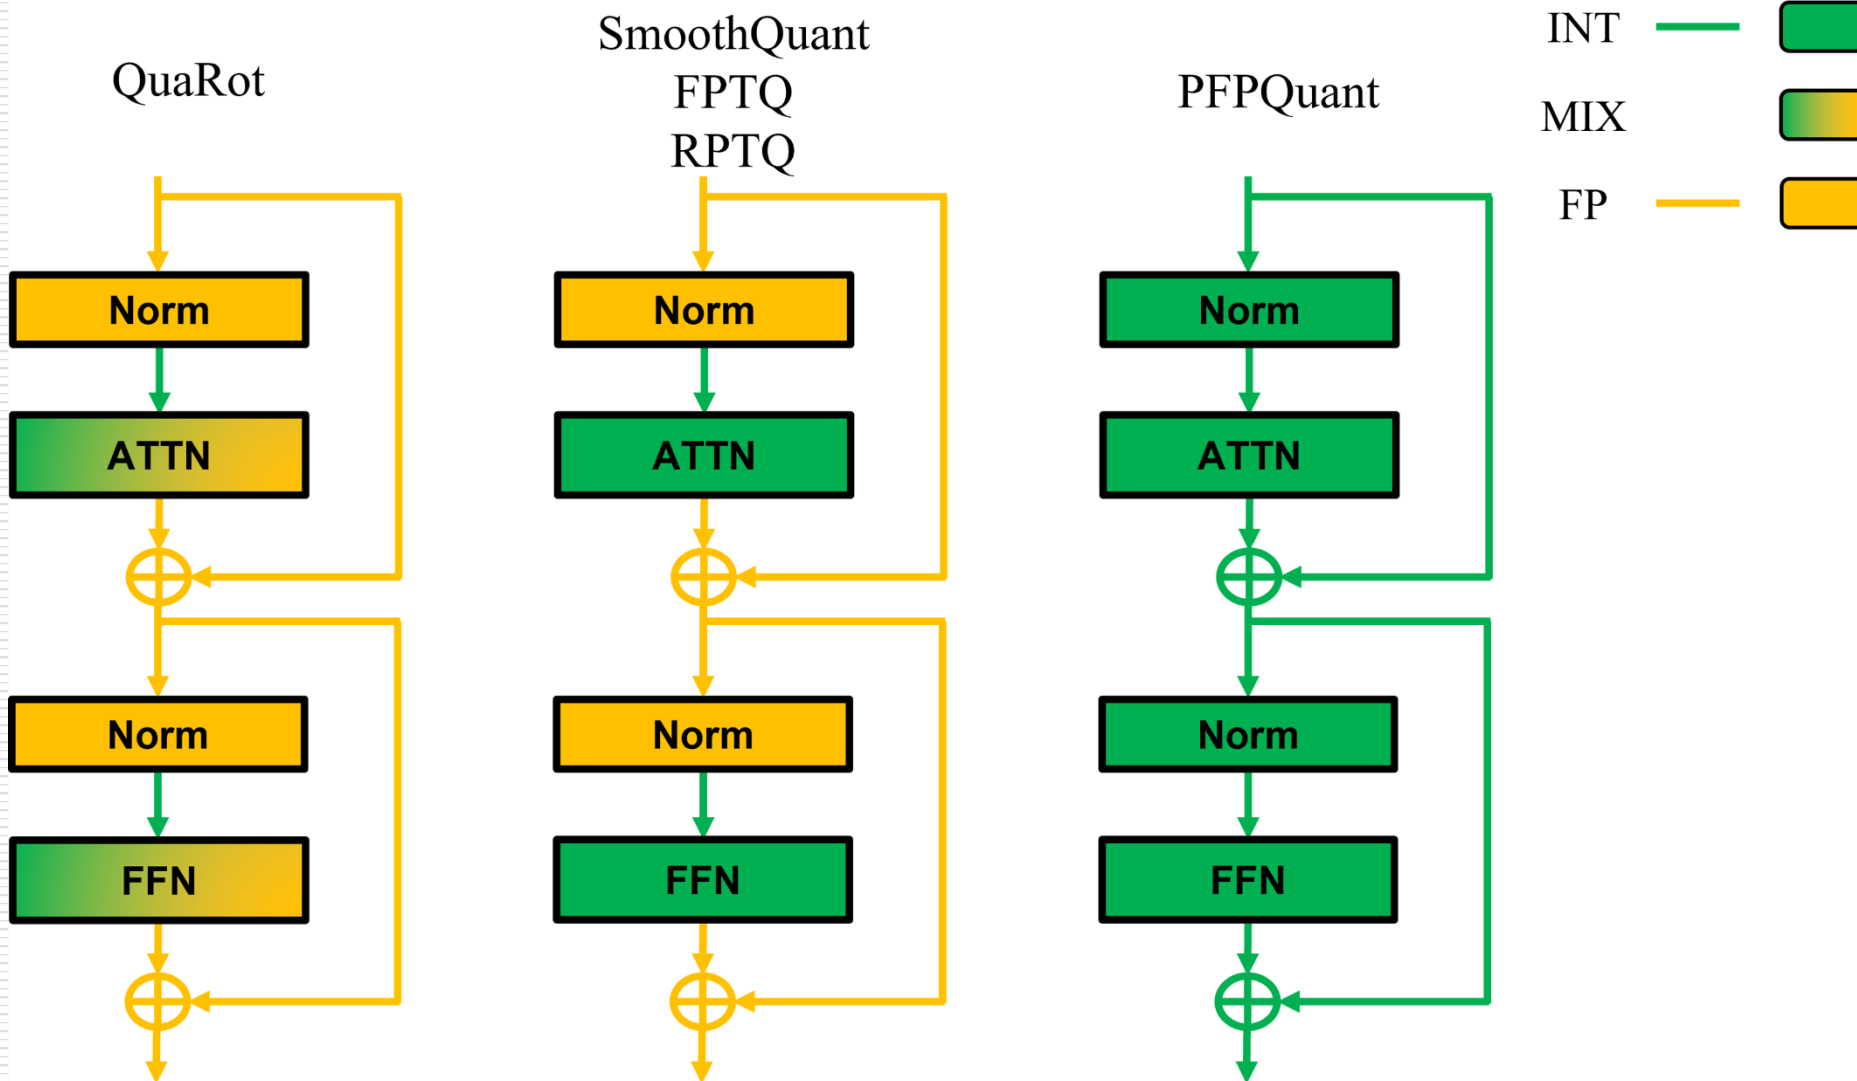
Table 2: **R** indicates the bit-width of RoPE sine and cosine tables, and **S** indicates the softmax output bit-width. For example, **R8-S8** uses 8-bit RoPE tables and clamps the softmax output to unsigned 8-bit (Q0.8). The data shows results after applying GOM, ORI-LN, and RL-SAT; omitting any of these strategies leads to pure fixed-point quantization failure.

Implementation Overview

- LLaMA-2 7B



Comparison



Homework

HW1 & HW2

- HW1 (35%)
 - 詳細介紹k-means分群演算法，並說明你覺得該演算法要怎麼用來壓縮NN模型
 - › 請善用圖示說明：圖最好自己畫，少用截圖
 - › 以簡報形式呈現
- HW2 (65%)
 - 請從下列論文擇一閱讀，並製作圖文並茂的簡報，圖例為主，文字需簡潔扼要：圖最好自己畫，少用截圖
 - › [AWQ: Activation-aware Weight Quantization for On-Device LLM Compression and Acceleration](#)
 - › [QuaRot: Outlier-free 4-bit Inference in Rotated LLMs](#)
 - › [A White Paper on Neural Network Quantization](#)
 - › [Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference](#)

Format

- 將HW1 & HW2**融合成一份**PowerPoint檔案(**不超過25頁投影片**)，繳交該檔案，檔名為：
 - 中文姓名_Lab12
- 將該份PowerPoint檔案轉成PDF檔案，一併繳交，檔名為：
 - 中文姓名_Lab12
- 總共要繳交**兩份**檔案，注意，**不要壓縮**
- 別直接傳兩個檔案，請至少附上姓名、主旨&禮貌
- 繳交至以下信箱：
 - anson.twhu.ee11@nycu.edu.tw

Deadline

- 23:59, Aug 26, 2024

Thank you