

AI training course

HW12

STuser19 賴昱凱

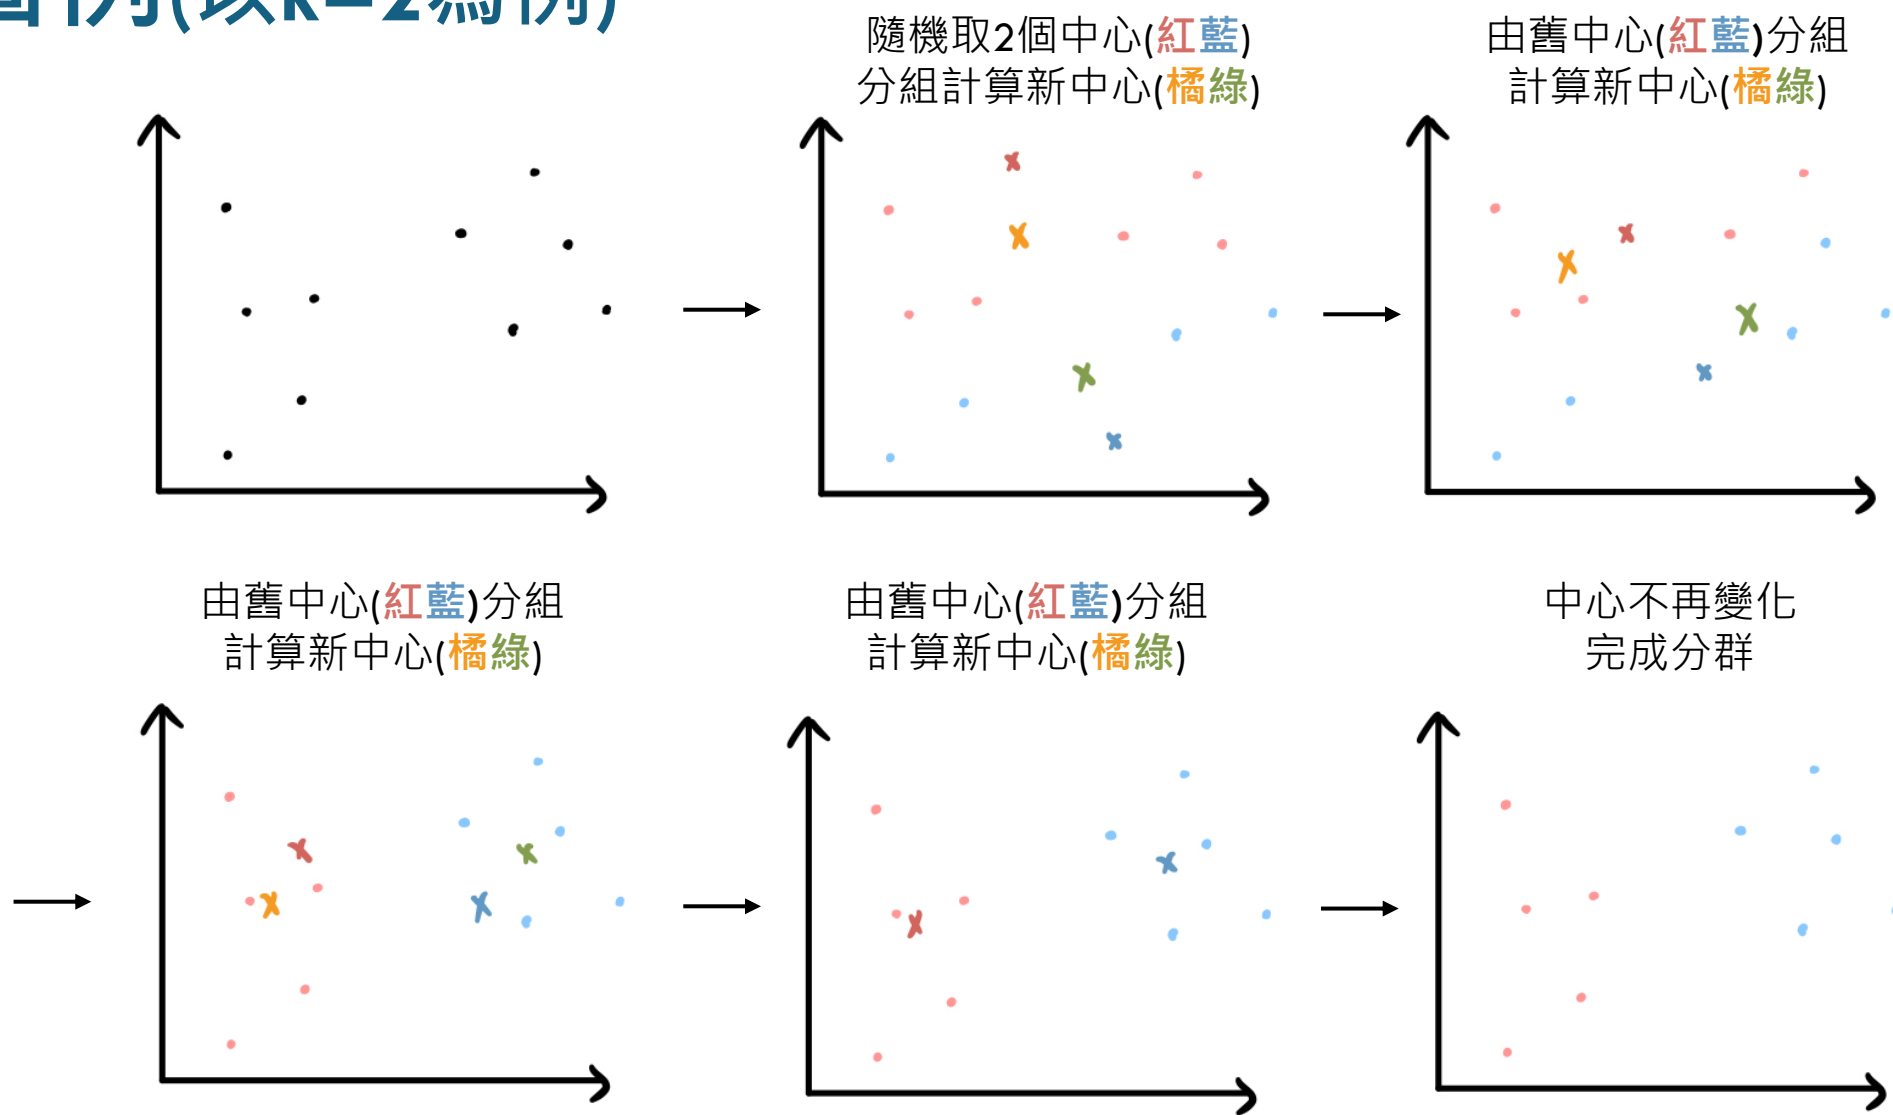
k-means分群演算法

將資料依照相似度分成k個群組

演算法

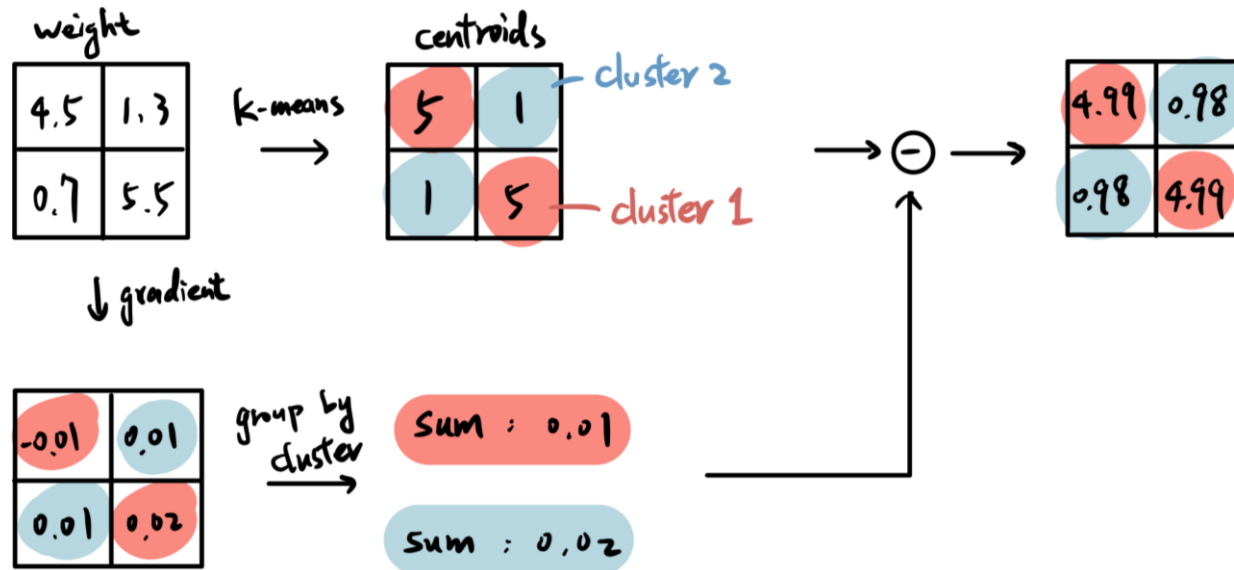
1. 隨機取 k 點作為中心
2. 計算各資料點與各中心的直線距離
3. 將資料點納入直線距離最近的中心群組
4. 計算群組的新中心
5. 重複執行3、4直到中心不再變化

圖例(以 $k=2$ 為例)




怎麼用來壓縮NN模型(一)

1. 利用k-means將權重分群，形成k個clusters
2. 每個Cluster中的資料皆由中心取代(可減少記憶體需求)
3. Fine-tuning with summation of corresponding gradients



怎麼用來壓縮NN模型(二)

1. 將data利用k-means分成k個clusters
2. Reorder by clusters
3. 各clusters分別做Quantization (可降低模型壓縮後的表現缺失)



Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference

Benoit Jacob, Skirmantas Kligys, Bo Chen, Menglong Zhu, Matthew Tang ,Andrew Howard, Hartwig Adam, Dmitry Kalenichenko

Google Inc.

Introduction

現狀問題

- CNN模型過大無法於mobile device使用

現有解法

- 新架構: MobileNet、SqueezeNet、ShuffleNet、DenseNet等
- Quantization: 用低位元降低儲存成本

現有Quantization缺陷

- 過度參數化的模型易壓縮，較無參考價值: AlexNet等
- 未於硬體驗證效果，且主要著重於減少儲存而非增加速度

Specific Contributions

Quantization scheme

- weight 、 activation: 8-bit int, Bias: 32-bit int

Quantized inference framework

- 於int-only hardware實現高效運算

Quantized training framework

- low-precision fixed-point arithmetic

benchmark results

- 基於MobileNet在ARM CPU執行，ImageNet及COCO表現優異平衡

Quantized Scheme

Inference : integer-only, training : floating-point

$$\textit{Quantization scheme} : r = S(q - Z) \quad (1)$$

r : real number

q : integer (quantized values)

S : positive real number (const)

Z : integer corresponding to $r = 0$ (zero-point, const)

Integer-arithmetic-only matrix multiplication

multiplication of two $N \times N$ matrices of real numbers r_1, r_2 and $r_3 = r_1 r_2$

$$r_{\alpha}^{i,j} = S_{\alpha} \left(q_{\alpha}^{i,j} - Z_{\alpha} \right), \quad \begin{cases} 1 \leq i, j \leq N \\ \alpha = 1, 2, 3 \end{cases} \quad (2)$$

From the definition of matrix multiplication

$$S_3(q_3^{i,k} - Z_3) = \sum_{j=1}^N \left(q_1^{i,j} - Z_1 \right) \left(q_2^{j,k} - Z_2 \right) \quad (3)$$

$$q_3^{i,k} = Z_3 + M \sum_{j=1}^N (q_1^{i,j} - Z_1)(q_2^{j,k} - Z_2) \quad (4), \quad M := \frac{S_1 S_2}{S_3} \quad (5)$$

We empirically find M to always be in the interval $(0, 1)$

$$\textit{normalized form} : M = 2^{-n}M_0, \quad \begin{cases} M_0 \in [0.5, 1) \\ n \in N \end{cases} \quad (6)$$

M_0 : *fixed – point multiplier*

- 以int32為例，表示 M_0 的整數為最接近 $2^{31}M_0$ 的int32值，且因 $M_0 \geq 0.5$ 故該值至少為 2^{30} 。
- 乘以 2^{-n} 可以用bit-shifting實現

Efficient handling of zero-points

Efficiently implement the evaluation of Equation (4), $O(2N^3)$

$$q_3^{i,k} = Z_3 + M \left(NZ_1Z_2 - Z_1a_2^k - Z_2\bar{a}_1^i + \sum_{j=1}^N q_1^{i,j}q_2^{j,k} \right) \quad (7)$$

$$\text{where } a_2^k = \sum_{j=1}^N q_2^{j,k}, \bar{a}_1^i = \sum_{j=1}^N q_1^{i,j} \quad (8), \text{ each } N \text{ additions}$$

a_2^k and \bar{a}_1^i : $2N^2$ additions, $\sum_{j=1}^N q_1^{i,j}q_2^{j,k}$ (9): $2N^3$ arithmetic operations

Implementation of a typical fused layer

- 將bias-addition 、 activation 融合至矩陣乘法
- Weights and activations use uint8 or int8
- Accumulator uses int32 to store (uint8 * uint8)
- Bias vector is quantized to int32 with $S_{\text{bias}} = S_1 S_2$, $Z=0$
- Output: int32 \rightarrow int8 [0, 225] (fixed-point multiplication)

Training with simulated quantization

Inference with int 、training with FP 的缺點

- 小模型表現差
- 數值範圍小的weight誤差大
- 異常值使精度降低

解法

在forward中simulates quantization，並維持bias、weight用FP

- Weights are quantized before ConV.
- Activations are quantized at points they would be during inference

Learning quantization ranges

Weight Quantization

- 範圍由 $[\min w, \max w]$ 量化為 $[-127, 127]$

Activation Quantization

- 由Training及EMA找範圍 $[a, b]$ 。訓練初期不使用防止遺失重要數值

Boundary Adjustment

- 微調 $[a, b]$ 確保 0.0 在量化後精確表示為integer

Quantization in TensorFlow

- training使用simulated quantization，inference使用low-bit quantization

Batch normalization folding

- 使用Batch Norm的模型

- Training : 獨立操作
- Inference : “folding” into layers. e.g. ConV. 包含Weight_{fold}及bias

$$w_{fold} := \frac{\gamma w}{\sqrt{\sigma_B^2 + \varepsilon}}, \left\{ \begin{array}{l} \gamma: \text{scale parameter} \\ \sigma_B^2: \text{estimated variance of ConV. results} \\ \varepsilon: \text{small constant} \end{array} \right.$$

EXP-Quantized training of Large Networks

Accuracy :

- for ResNet and InceptionV3, FP vs. int accuracy is within 2%

Different Quantization

- ResNet

Integer – only \approx 5 – bit INQ > 2 – bit FGQ

- Inception v3

7 – bit \approx 8 – bit

ReLU6表現較佳

EXP-Quantization of MobileNets

MobileNet Quantization (within the same runtime)

- accuracy : Integer-only quantized > floating-point models for $\sim 10\%$
- Snapdragon 835 LITTLE cores : 33ms latency for real-time (30 fps) operation

COCO Object Detection

- 2x speedup with only $\sim 1.8\%$ accuracy loss

Face Detection and Attribute Classification

- reduce latency by $\sim 2x$ with only $\sim 2\%$ accuracy loss
- Single, Multi core 皆有大幅加速